

# Operating Systems: Deadlocks

CSC-4320/6320 –Summer 2014

# Real-life Deadlock



- Three kinds of OS-deadlock solutions:
  1. have mechanisms so that a deadlock never happens in the first place
  2. detect that we're in a deadlock, and do something to fix it
  3. do nothing and when things don't work have the "operator" reboot it all

# Deadlocks

- 20<sup>th</sup> century Kansas Law: “When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone”
- Deadlock with two threads and two resources (see Figure 7.4: Java example)

Thread #1	Thread #2
lock(A)	lock(B)
lock(B)	lock(A)
unlock(B)	unlock(A)
unlock(A)	unlock(B)

- Typically it is the responsibility of the programmer to avoid deadlocks
  - Deadlocks should be rare if the burden's placed on programmers who are highly motivated to avoid deadlocks
  - A manual restart (i.e., kill-restart) is always an option
  - Therefore, avoid making the OS more complicated and let users fend for themselves
  - e.g., Windows/Linux provide no help in this matter
- We're going to look at what OSes could provide, because understanding this leads us to understanding how to avoid deadlocks in our own programs

# System Model

- System consists of
  - some **resources**
  - **resource types**:  $R_1, R_2, \dots, R_m$ 
    - There could be one or more resource in each type
    - e.g., Physical: 4 printers, 2 network cards
      - each protected by an associated lock
      - either visible to the application, or within the Kernel
      - e.g., when you do an `open()`, there is a lock in the Kernel for that file
    - **processes**:  $P_1, P_2, \dots, P_n$
  - Each process can:
    - request a resource of a given type
      - And block/wait until one resource instance of that type becomes available
    - use a resource
    - release a resource

# Deadlock State

- We have a deadlock if every process  $P_i$  is waiting for a resource instance that is being held by another process
- A deadlock can arise only if all four conditions hold
  - **Mutual Exclusion**
    - At least one resource is non-sharable: at most one process at a time can use it
  - **Hold-and-Wait**
    - At least one process is holding one resource while waiting to acquire others, that are being held by other processes
  - **No preemption**
    - A resource cannot be preempted (a process needs to give it up voluntarily)
  - **Circular Wait**
    - There exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that
      - $P_l$  is waiting for a resource that is held by  $P_{l+1}$
      - $P_n$  is waiting for a resource that is held by  $P_0$

# Deadlock State


- The four conditions:
  - Mutual Exclusion
  - Hold-and-Wait
  - No preemption
  - Circular Wait
- Note that “circular wait” implies “Hold-and-Wait”
  - It is useful to separate them, as we’ll see later
- The four conditions together are only a **necessary condition**
  - If the four conditions hold, there **may** be a deadlock
  - If there is a deadlock, then the four conditions hold

# Resource Allocation Graphs

- Describing the system can be done precisely and easily with a system resource-allocation graph
- The graph contains:
  - A set of vertices,  $V$ , that contains
    - One vertex for each process:  $\{P_0, P_1, \dots, P_n\}$
    - One vertex for each resource type:  $\{R_1, R_2, \dots, R_m\}$ 
      - Which indicates the number of resource instances for that type



vertex for process  $P_i$



vertex for resource type  $R_j$   
with 3 resource instances

# Resource Allocation Graphs

- The graph contains:
  - A set of directed edges,  $E$ , that contains
    - Request edge: from  $P_i$  to  $R_j$  if process  $P_i$  has requested a resource of type  $R_j$ 
      - Points to the resource type rectangle
    - Assignment edge: from  $R_j$  instance to process  $P_i$  if  $P_i$  holds a resource instance of type  $R_j$ 
      - Points from a dot inside the resource type rectangle




- If a resource request can be fulfilled, then a request edge is transformed into an assignment edge
- When a process releases a resource, the assignment edge is deleted



# Example Resource Graph

- Figure 7.1




# Graphs and Deadlocks

- Theorem:
  - If the graph contains no (directed) cycles, then there is no deadlock
  - If the graph contains a cycle, then there may be a deadlock
- If there is only one resource instance per resource type, then we have a stronger Theorem:
  - The existence of a cycle is a sufficient and necessary condition for the existence of a deadlock
    - Each process involved in the cycle is deadlocked


# Cycle and Deadlock

- Figure 7.2



# Cycles and No Deadlock


- Figure 7.3



# Simple Example

- 8 resources
- 2 threads
- Each thread does:

```
while(true){  
    for(int i=0; i<M; i++)  
        <grab a resource>  
    <do some work>  
    for(int i=0; i<M; i++)  
        <release a resource>  
}
```



**Question:** What is the largest  $M$  value to guarantee no deadlocks?

# Deadlock Handling

- What do we do about deadlocks?
  - We can prevent deadlocks
    - Deadlock prevention
      - Ensure that one of the four conditions never holds
    - Deadlock avoidance
      - Use information about future resource usage of processes
  - We can identify deadlocks and take action
    - Deadlock detection and recovery
      - An algorithm for deadlock detection
      - A recovery strategy
    - We can do nothing and hope
      - That is what Windows, Linux, and the JVM do
      - Eventually the deadlock may snowball until the system no longer functions and requires a manual intervention (restart)

# Deadlock Prevention

- The four conditions:
  - Mutual Exclusion
  - Hold-and-Wait
  - No preemption
  - Circular Wait
- Getting rid of Mutual Exclusion?
  - IN general we cannot design a system in which we do not have some type of exclusion on some types of resources
- Getting rid of No Preemption?
  - This would force resource release from a waiting process (A) that holds a resource needed by another process (B)
  - A is restarted later and must reacquire all its resources
  - This is easily done for resources that have an easily saved/restored state (e.g., CPU with registers)
  - But cannot be done in general as the processes may be in the middle of doing something that leaves an inconsistent state

# Getting Rid of Hold and Wait

- A process cannot request a resource if it holds any other resource
- **Option #1:** a process could be allocated all the resources it needs before it begins execution
  - Problem: low resource utilization
    - A resource is held during the whole process lifetime even if it is used for a tiny fraction of it
- **Option #2:** a process can request a (bulk of) resource(s) only if it holds no other resources
  - Problem: may not be possible to implement every process as a sequence of “acquire N/release N” steps
- Problem in both options: starvation is possible
  - Some other process may always hold one of the needed resources and acquiring them one after the other is the only way



# Getting Rid of Circular Wait

- Preventing cycles:
  - Impose a total ordering on resource types
    - An integer value is assigned to each type
  - A process must request resources by increasing type order
    - or, must release all resources of higher order before requesting a resource of lower order
- The above will prevent circular wait
  - Simple proof by contradiction in Section 7.4.4
- This works trivially for the two-lock deadlock ( $A < B$ )

Thread #1

lock(A)

lock(B)

unlock(B)

unlock(A)

Thread #2

lock(B)

~~lock(A)~~ illegal

unlock(A)

unlock(B)

# Getting Rid of Circular Wait

- It is up to application developers to follow the order
  - Otherwise code will simply say “fail”
- It may not be easy to define the order a-priori
  - If some process may need resource type A before type B, and some other may need resource type B before type A, then you cannot define the order
  - Hard to figure out an order for all system resources
- FreeBSD provides an order-verifier for locks
  - It records lock usage order
  - And then later enforces the recorded order
  - Pretty simple to implement

# Deadlock Avoidance


- Idea: if I know what resources a process will need in the future, perhaps I can anticipate deadlocks
- A simple and useful model: each process declares the maximum number of resource of each type that it may need
- Resource state:
  - The number of available resources in each type
  - The number of assigned resources in each type
  - The maximum number of resources of each type for process
- Goal: ensure that we are always in a **safe state**

# Safe State

- Definition of a **safe state**: there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the system such that
  - for each  $P_i$ , the resource that  $P_i$  can still request can be satisfied by currently available resources + resources held by all the  $P_j$ , with  $j < i$
  - That is (for  $j < i$ ):
    - If  $P_i$  resource needs are not immediately available, then  $P_i$  must wait until all  $P_j$  have finished
    - When each  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and eventually terminate
    - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources, and so on...
  - Such a sequence is called a **safe sequence**
- A state without a safe sequence is called **unsafe**

# Safe State

- Theorem:
  - If there is a deadlock, then the state is unsafe
  - If the state is unsafe, then there may be a deadlock



- Goal: never enter an unsafe state, period
  - And conservatively precluded non-deadlocked unsafe states

# Example from Section 7.5.1

- 12 Resources of the same type, 9 assigned
- 3 processes:


	<u>Maximum Need</u>	<u>Currently Holds</u>
$P_0$	10	5
$P_1$	4	2
$P_2$	9	2

- A safe sequence:  $\langle P_1, P_0, P_2 \rangle$
- If one gives 1 resource to  $P_2$ , then we get to an unsafe state
  - $P_2$  holds 3,  $P_1$  gets and releases 2, then neither  $P_0$  nor  $P_2$  can get everything they need
- Looking at state safety could be done using a brute-force (high-complexity) algorithm

# Graph-based Avoidance Alg.

- If each resource type has only one instance, then it is easy to avoid deadlocks
  - A more complex algorithm called the “Banker’s algorithm” must be used for multiple instances
- Build a resource allocation graph, but add **claim edges**
  - Edges that correspond to potential future resource needs (all of them)
    - Depicted with a dashed line
  - When a resource is assigned, replace the claim edge with an assignment edge
- Grant resource allocation only if it does not create a cycle in the resource allocation graph
  - The cycle may contain claim edges
  - Detecting a cycle in a graph with  $n$  vertices can be done in  $n^2$  time

# Graph Example



- There is a cycle in the graph
- The request is denied
  - It could lead to a deadlock




# Deadlock Detection-Recovery


- Detection-Recovery:
  - Allows system to enter a deadlock state
  - Detect the deadlock state
  - Take some appropriate action to recover
- In the case of one instance per resource type, detection is simple
  - Build the resource allocation graph
  - Run an  $O(n^2)$  cycle-detection algorithm
- Otherwise a more complex algorithm is needed
  - Uses ideas from the Banker's algorithm (see Section 7.5.3 if interested)

# Deadlock Detection Example

- Figure 7.11



(a)



(b)

# Deadlock Detection-Recovery

- How often should one run the detection algorithm?
  - Run it often: expensive, but good if deadlocks are frequent
    - extreme: for each resource request, in which case one knows which process “caused” the deadlock
  - Run it rarely: cheap, but bad if deadlocks are frequent
    - and it will be difficult to tell which process “caused” the deadlock

# Deadlock Detection-Recovery

- What about recovery?
- Two kinds of actions
  - Process termination
  - Resource preemption
- Process termination
  - Kill all deadlocked processes
    - may be wasteful
  - Kill one process at a time until the deadlock disappears
    - High overhead because deadlock detection algorithm is run at each step (but the system was frozen anyway)
  - Killing a process could be tricky
    - The process may be in the middle of something that would leave an inconsistent state, that must be fixed

# Deadlock Detection-Recovery

- Resource Preemption
  - Selecting a victim: which resource/process need to be preempted
  - Rollback: when preempting a resource from a process, that process must be rolled back
    - Simple solution: restart the process from scratch
    - May require inconsistent state cleanup
  - Starvation: ensure that one process does not see its resources preempted from it forever

# Conclusion

- Three methods
  1. Deadlock prevention/avoidance
  2. Deadlock detection-recovery
  3. Do nothing and let users deal with it
- The solutions we have discussed for 1 and 2 above are interesting
- Most argue that none of them covers all the bases
- One could combine them all and be effective, but at the cost of much increased Kernel complexity
- Therefore, in practice, it is option 3 that gets used