

Operating Systems: Midterm Review

CSC-4320/6320 –Summer 2014

What to Study?

- Slides, notes, and reading assignments up to and including “Deadlocks”
- Past homework assignments
- Past programming assignments
- There should be very few surprises
- Type of Questions:
 - general “what do you think?” questions
 - homework-like problems
- Exam will be available for 24 hours starting 6:00 p.m. Thursday.
 - You are expected to work alone
 - You may use your book, notes, google
 - I am more interested in seeing that you can find the answers and have an understanding of the topic than I am in your ability to memorize.

Sample General Questions

- What is the difference between an Interrupt and an Exception/Trap? Give two examples of each

Sample General Questions

- What is the difference between an Interrupt and an Exception/Trap? Give two examples of each
 - Both are events that the OS must react to
 - An Interrupt is due to some external event
 - e.g., keyboard input
 - e.g., disk operation completion
 - A trap/exception is caused by an instruction
 - e.g., illegal memory access
 - e.g., system call instruction

Sample General Question

- What are the advantages/disadvantages of User vs. Kernel Threads?

Sample General Question

- What are the advantages/disadvantages of User vs. Kernel Threads?
 - User-level
 - good: very low overhead for creation/maintenance because the kernel's not involved
 - bad: cannot take advantage of multi-core architectures
 - bad: if one thread blocks, they all block
 - Kernel-level
 - bad: high overhead because the kernel is involved
 - removes the two “bad” of user threads

Sample General Question

- Explain what happens during a context-switch

Sample General Question

- Explain what happens during a context-switch
 - Assume A is running and we are switching to B because A's time quantum has expired
 - Save A's "state" (registers) into its PCB structure in the kernel
 - Move A's PCB to the "Ready" queue
 - Restore B's "state" (registers) into the CPU by reading it from B's PCB
 - Move B's PCB to the "Running" queue
 - Tell the CPU to start a new fetch-decode-execute cycle
 - The PC register contains the address of the next instruction that B must execute

Sample General Questions

- What is a zombie?
- What is an orphan?

Sample General Questions

- What is a zombie?
 - A process that has terminated (i.e., called “exit()”), but whose parent has not acknowledged the death (i.e., not called wait() or waitpid()).
 - It is kept around so that later the parent can find out its return value (e.g., exit(42));
- What is an orphan?
 - A process whose parent has died
 - It is adopted by process PID=1 (meaning that its PPID=1)

Sample General Question

- What is priority inversion?
- Why should context-switching overhead be low?
- What happens if at the end of a time quantum of a running process the ready queue is empty?
- etc. etc.

Sample General Questions

- There may be a general question about the programming assignments
- Just to make sure that you've done them and understood concepts in there
- Example: What system program is used to trace system calls on a Linux system? **strace**
- Example: What can be done to force the init process to adopt a child process you create?
Create a child process that creates a child and then calls exit, thus the grandchild will be adopted by init.

Preparing for General Questions

- Most of the lectures consist of topics that beg for a “how does this work?” question
- Each such topic corresponds to an OS concept
 - There are not that many such concepts in each lecture
- Go through general questions in homework assignments, questions in the chapter practice exercises may be a good place to review your understanding of the chapter

Problems

- All problems will resemble homework problems with maybe a few more creative questions added
 - Given a program with calls to `fork()`, what happens?
 - Given a set of jobs with arrival time, burst time, and priority, what do various scheduling algorithms do?
 - Given a multi-threaded program, what are the possible outputs?
 - Given a multi-threaded program, how would you add locks?
 - Given some scenario, is there a possible deadlock?

Problems

- Some problems will be: “look at that code, and tell me what it does”
 - either pseudo-code or C++/Java code

Concurrency

- Since we didn't have a concurrency Homework assignment, we can do a few "similar" concurrency examples..
- Spinning locks vs. blocking locks
- Interleaved output
- Locks
- Deadlocks

Spinning vs. Blocking Locks

- Spinning:
 - burn CPU cycles checking the lock continuously, which is wasteful
 - But as soon as the lock is released by whoever had it you grab it, which is good
- Blocking lock:
 - Go to “sleep” asking for somebody to wake up when the key’s ready, which avoids being a useless CPU hog
 - But this requires much more work as the OS is now involved to put you to sleep and wake you up
 - Moving your PCB from various queues, etc.
- Rules of thumb:
 - Spinning for a long time is wasteful
 - Blocking if the lock is never taken for a long time is wasteful

Interleaved Output

- Two threads, two global variables $a=1$ and $b=1$
- Thread #1: $b++$; $a=b-a$
- Thread #2: $a++$
- What are the possible outputs?
- 3 “clean” interleavings:
 - $a++$; $b++$; $a=b-a$ (0,2)
 - $b++$; $a++$; $a=b-a$ (0,2)
 - $b++$; $a=b-a$; $a++$ (2,2)
- There can also be bad interleaving of the non-atomic “ $a++$ ” and “ $a=b-a$ ” “instructions” since they both write to the same variable
- These can happen only after $b++$ has executed and then we’re left with both “ $a++$ ” and “ $a=b-a$ ” to execute

Interleaved Output

- What are the “bad” interleavings for “a++” and “a=b-a”?
 - Remember that at this point $b=2$
- **Case #1:**
 - “a++” load the value of a, which is 1, computes 2, and is about to write back to a, but there is a context switch
 - “a=b-a” loads the value of b (2) and of a(1), computes $b-a=1$, and is about to write back to a, but there is a context switch
 - “a++” resumes and writes back value 2 to a
 - “a=b-a” resumes and overwrites value 2 by value 1
 - So in the end we have (1,2)

Interleaved Output

- What are the “bad” interleavings for “a++” and “a=b-a”?
 - Remember that at this point $b=2$
- Case #2
 - “a=b-a” loads the value of $b(2)$ and of $a(1)$, computes $b-a=1$, and is about to write back to a , but there is a context switch
 - “a++” loads the value of a , which is 1, computes 2, and is about to write back to a , but there is a context switch
 - “a=b-a” resumes and writes back value 1 to a
 - “a++” resumes and overwrites value 1 by value 2
 - So in the end we have (2,2)

Interleaved Output

- In the end we found these possible outputs:
 - Clean #1: (0,2)
 - Clean #2: (0,2)
 - Clean #3: (2,2)
 - Unclean #1: (1,2)
 - Unclean #2: (2,2)
- So, in total, we have 3 possible outputs: (0,2), (2,2), and (1,2)

Locks

- When writing concurrent code in which multiple threads update the same (global) variables, one must make sure to avoid all “unclean” interleaving
- For this purpose we create critical sections
- Each critical section (which can have multiple zones in the code!) is defined by a lock, with a `lock(A)/unlock(A)` pair of statements

Locks

- Say our code has the following statements and that many threads will execute these statements

....

x++;

...

y++;

x++;

...

z=2;

...

Locks

- Say our code has the following statements and that many threads will execute these statements

```
....  
x++;  
...  
y++;  
x++;  
...  
z=2;  
...
```
- A brute force approach is to use one lock

```
...  
lock(A);  
x++;  
unlock(A);  
...  
lock(A);  
y++;  
x++;  
unlock(A);  
...  
lock(A);  
z=2;  
unlock(A);  
...
```


Locks

- Say our code has the following statements and that many threads will execute these statements
- A brute force approach is to use one lock

```
...  
x++;  
...  
y++;  
x++;  
...  
z=2;  
...
```

```
...  
lock(A);  
x++;
```

NOT good because a thread cannot do x++ while another one is doing y++.
Loss of concurrency, and thus performance.
If you want everything to be sequential, why use threads in the first place

```
...  
unlock(A);  
...  
lock(A);  
z=2;  
unlock(A);  
...
```

Locks

- Say our code has the following statements and that many threads will execute these statements

```
....  
x++;  
...  
y++;  
x++;  
...  
z=2;  
...
```

- First Question to answer, how many locks?
- Should there be a lock for protecting updates of x?
- Should there be a lock for protecting updates of y?
- Should there be a lock for protecting updates of z?

Locks

- Say our code has the following statements and that many threads will execute these statements

```
....  
x++;  
...  
y++;  
x++;  
...  
z=2;  
...
```

- Solution:

```
...  
lock(A);  
x++;  
unlock(A);  
...  
lock(B);  
y++;  
unlock(B);  
lock(A);  
x++;  
unlock(A);  
...  
lock(A);  
z=2;  
unlock(A);  
...
```

Deadlocks

- We haven't seen much on this topic
- Make sure you understand the resource allocation graphs
 - If the “gray boxes” have more than “one dot”: if there is a cycle, there may be a deadlock
 - If the “gray boxes” have only “one dot”: if there is a cycle, there is a deadlock